

# workshop2-ai-ml

March 11, 2025

## 0.1 Some Helper Function:

### 0.1.1 Softmax Function:

```
[1]: import numpy as np

def softmax(z):
    """
    Compute the softmax probabilities for a given input matrix.

    Parameters:
    z (numpy.ndarray): Logits (raw scores) of shape (m, n), where
        - m is the number of samples.
        - n is the number of classes.

    Returns:
    numpy.ndarray: Softmax probability matrix of shape (m, n), where
        each row sums to 1 and represents the probability
        distribution over classes.

    Notes:
    - The input to softmax is typically computed as:  $z = XW + b$ .
    - Uses numerical stabilization by subtracting the max value per row.
    """

    z_max = np.max(z, axis=1, keepdims=True)
    exp_z = np.exp(z - z_max)
    softmax_probs = exp_z / np.sum(exp_z, axis=1, keepdims=True)

    return softmax_probs
```

### 0.1.2 Softmax Test Case:

This test case checks that each row in the resulting softmax probabilities sums to 1, which is the fundamental property of softmax.

```
[2]: # Example test case
z_test = np.array([[2.0, 1.0, 0.1], [1.0, 1.0, 1.0]])
softmax_output = softmax(z_test)
```

```

# Verify if the sum of probabilities for each row is 1 using assert
row_sums = np.sum(softmax_output, axis=1)

# Assert that the sum of each row is 1
assert np.allclose(row_sums, 1), f"Test failed: Row sums are {row_sums}"

print("Softmax function passed the test case!")

```

Softmax function passed the test case!

### 0.1.3 Prediction Function:

```

[3]: def predict_softmax(X, W, b):
    """
    Predict the class labels for a set of samples using the trained softmax
    ↪ model.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of
    ↪ samples and d is the number of features.
    W (numpy.ndarray): Weight matrix of shape (d, c), where c is the number of
    ↪ classes.
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    numpy.ndarray: Predicted class labels of shape (n,), where each value is
    ↪ the index of the predicted class.
    """
    logits = np.dot(X, W) + b
    probabilities = softmax(logits)
    predicted_classes = np.argmax(probabilities, axis=1)

    return predicted_classes

```

### 0.1.4 Test Function for Prediction Function:

The test function ensures that the predicted class labels have the same number of elements as the input samples, verifying that the model produces a valid output shape.

```

[4]: # Define test case
X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3
↪ samples, 2 features)
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3
↪ classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)

```

```

# Expected Output:
# The function should return an array with class labels (0, 1, or 2)

y_pred_test = predict_softmax(X_test, W_test, b_test)

# Validate output shape
assert y_pred_test.shape == (3,), f"Test failed: Expected shape (3,), got {y_pred_test.shape}"

# Print the predicted labels
print("Predicted class labels:", y_pred_test)

```

Predicted class labels: [1 1 0]

### 0.1.5 Loss Function:

```

[5]: def loss_softmax(y_pred, y):
    """
    Compute the cross-entropy loss for a single sample.

    Parameters:
    y_pred (numpy.ndarray): Predicted probabilities of shape (c,) for a single
    sample,
                                where c is the number of classes.
    y (numpy.ndarray): True labels (one-hot encoded) of shape (c,), where c is
    the number of classes.

    Returns:
    float: Cross-entropy loss for the given sample.
    """

    epsilon = 1e-12
    y_pred = np.clip(y_pred, epsilon, 1.0 - epsilon)
    loss = -np.sum(y * np.log(y_pred))

    return loss

```

### 0.2 Test case for Loss Function:

This test case Compares loss for correct vs. incorrect predictions. \* Expects low loss for correct predictions. \* Expects high loss for incorrect predictions.

```

[6]: import numpy as np

# Define correct predictions (low loss scenario)
y_true_correct = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True one-hot
labels

```

```

y_pred_correct = np.array([[0.9, 0.05, 0.05],
                           [0.1, 0.85, 0.05],
                           [0.05, 0.1, 0.85]]) # High confidence in the
↳ correct class

# Define incorrect predictions (high loss scenario)
y_pred_incorrect = np.array([[0.05, 0.05, 0.9], # Highly confident in the
↳ wrong class
                             [0.1, 0.05, 0.85],
                             [0.85, 0.1, 0.05]])

# Compute loss for both cases
loss_correct = loss_softmax(y_pred_correct, y_true_correct)
loss_incorrect = loss_softmax(y_pred_incorrect, y_true_correct)

# Validate that incorrect predictions lead to a higher loss
assert loss_correct < loss_incorrect, f"Test failed: Expected loss_correct <
↳ loss_incorrect, but got {loss_correct:.4f} >= {loss_incorrect:.4f}"

# Print results
print(f"Cross-Entropy Loss (Correct Predictions): {loss_correct:.4f}")
print(f"Cross-Entropy Loss (Incorrect Predictions): {loss_incorrect:.4f}")

```

Cross-Entropy Loss (Correct Predictions): 0.4304  
Cross-Entropy Loss (Incorrect Predictions): 8.9872

### 0.2.1 Cost Function:

```

[7]: def cost_softmax(X, y, W, b):
    """
    Compute the average softmax regression cost (cross-entropy loss) over all
    ↳ samples.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of
    ↳ samples and d is the number of features.
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c), where n
    ↳ is the number of samples and c is the number of classes.
    W (numpy.ndarray): Weight matrix of shape (d, c).
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    float: Average softmax cost (cross-entropy loss) over all samples.
    """

    logits = np.dot(X, W) + b
    probabilities = softmax(logits)

```

```

epsilon = 1e-12
probabilities = np.clip(probabilities, epsilon, 1.0 - epsilon)
total_loss = -np.sum(y * np.log(probabilities))
n = X.shape[0]

return total_loss / n

```

## 0.2.2 Test Case for Cost Function:

The test case assures that the cost for the incorrect prediction should be higher than for the correct prediction, confirming that the cost function behaves as expected.

```

[8]: import numpy as np

# Example 1: Correct Prediction (Closer predictions)
X_correct = np.array([[1.0, 0.0], [0.0, 1.0]]) # Feature matrix for correct
↳ predictions
y_correct = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded,
↳ matching predictions)
W_correct = np.array([[5.0, -2.0], [-3.0, 5.0]]) # Weights for correct
↳ prediction
b_correct = np.array([0.1, 0.1]) # Bias for correct prediction

# Example 2: Incorrect Prediction (Far off predictions)
X_incorrect = np.array([[0.1, 0.9], [0.8, 0.2]]) # Feature matrix for
↳ incorrect predictions
y_incorrect = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded,
↳ incorrect predictions)
W_incorrect = np.array([[0.1, 2.0], [1.5, 0.3]]) # Weights for incorrect
↳ prediction
b_incorrect = np.array([0.5, 0.6]) # Bias for incorrect prediction

# Compute cost for correct predictions
cost_correct = cost_softmax(X_correct, y_correct, W_correct, b_correct)

# Compute cost for incorrect predictions
cost_incorrect = cost_softmax(X_incorrect, y_incorrect, W_incorrect,
↳ b_incorrect)

# Check if the cost for incorrect predictions is greater than for correct
↳ predictions
assert cost_incorrect > cost_correct, f"Test failed: Incorrect cost
↳ {cost_incorrect} is not greater than correct cost {cost_correct}"

# Print the costs for verification
print("Cost for correct prediction:", cost_correct)

```

```
print("Cost for incorrect prediction:", cost_incorrect)

print("Test passed!")
```

Cost for correct prediction: 0.0006234364133349324  
 Cost for incorrect prediction: 0.29930861359446115  
 Test passed!

### 0.2.3 Computing Gradients:

```
[9]: def compute_gradient_softmax(X, y, W, b):
      """
      Compute the gradients of the cost function with respect to weights and
      ↪biases.

      Parameters:
      X (numpy.ndarray): Feature matrix of shape (n, d).
      y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
      W (numpy.ndarray): Weight matrix of shape (d, c).
      b (numpy.ndarray): Bias vector of shape (c,).

      Returns:
      tuple: Gradients with respect to weights (d, c) and biases (c,).
      """

      logits = np.dot(X, W) + b
      probabilities = softmax(logits)
      error = probabilities - y
      grad_W = np.dot(X.T, error) / X.shape[0]
      grad_b = np.sum(error, axis=0) / X.shape[0]

      return grad_W, grad_b
```

### 0.2.4 Test case for compute\_gradient function:

The test checks if the gradients from the function are close enough to the manually computed gradients using `np.allclose`, which accounts for potential floating-point discrepancies.

```
[10]: import numpy as np

      # Define a simple feature matrix and true labels
      X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3
      ↪samples, 2 features)
      y_test = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True labels (one-hot
      ↪encoded, 3 classes)

      # Define weight matrix and bias vector
```

```

W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3
↳classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)

# Compute the gradients using the function
grad_W, grad_b = compute_gradient_softmax(X_test, y_test, W_test, b_test)

# Manually compute the predicted probabilities (using softmax function)
z_test = np.dot(X_test, W_test) + b_test
y_pred_test = softmax(z_test)

# Compute the manually computed gradients
grad_W_manual = np.dot(X_test.T, (y_pred_test - y_test)) / X_test.shape[0]
grad_b_manual = np.sum(y_pred_test - y_test, axis=0) / X_test.shape[0]

# Assert that the gradients computed by the function match the manually
↳computed gradients
assert np.allclose(grad_W, grad_W_manual), f"Test failed: Gradients w.r.t. W
↳are not equal.\nExpected: {grad_W_manual}\nGot: {grad_W}"
assert np.allclose(grad_b, grad_b_manual), f"Test failed: Gradients w.r.t. b
↳are not equal.\nExpected: {grad_b_manual}\nGot: {grad_b}"

# Print the gradients for verification
print("Gradient w.r.t. W:", grad_W)
print("Gradient w.r.t. b:", grad_b)

print("Test passed!")

```

```

Gradient w.r.t. W: [[ 0.1031051  0.01805685 -0.12116196]
 [-0.13600547  0.00679023  0.12921524]]
Gradient w.r.t. b: [-0.03290036  0.02484708  0.00805328]
Test passed!

```

### 0.2.5 Implementing Gradient Descent:

```

[11]: def gradient_descent_softmax(X, y, W, b, alpha, n_iter, show_cost=False):
      """
      Perform gradient descent to optimize the weights and biases.

      Parameters:
      X (numpy.ndarray): Feature matrix of shape (n, d).
      y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
      W (numpy.ndarray): Weight matrix of shape (d, c).
      b (numpy.ndarray): Bias vector of shape (c,).
      alpha (float): Learning rate.
      n_iter (int): Number of iterations.
      show_cost (bool): Whether to display the cost at intervals.

```

```

Returns:
tuple: Optimized weights, biases, and cost history.
"""

cost_history = []

for i in range(n_iter):
    grad_W, grad_b = compute_gradient_softmax(X, y, W, b)

    W -= alpha * grad_W
    b -= alpha * grad_b

    cost = cost_softmax(X, y, W, b)
    cost_history.append(cost)

    if show_cost and (i % (n_iter // 10) == 0 or i == n_iter - 1):
        print(f"Iteration {i + 1}/{n_iter}, Cost: {cost:.6f}")

return W, b, cost_history

```

### 0.3 Preparing Dataset:

```

[12]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

def load_and_prepare_mnist(csv_file, test_size=0.2, random_state=42):
    """
    Reads the MNIST CSV file, splits data into train/test sets, and plots one
    image per class.

    Arguments:
    csv_file (str) : Path to the CSV file containing MNIST data.
    test_size (float) : Proportion of the data to use as the test set
    (default: 0.2).
    random_state (int) : Random seed for reproducibility (default: 42).

    Returns:
    X_train, X_test, y_train, y_test : Split dataset.
    """

    # Load dataset
    df = pd.read_csv(csv_file)

    # Separate labels and features

```



```

y = df.iloc[:, 0].values # First column is the label
X = df.iloc[:, 1:].values # Remaining columns are pixel values

# Normalize pixel values (optional but recommended)
X = X / 255.0 # Scale values between 0 and 1

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
↳test_size=test_size, random_state=random_state)

# Plot one sample image per class
plot_sample_images(X, y)

return X_train, X_test, y_train, y_test

def plot_sample_images(X, y):
    """
    Plots one sample image for each digit class (0-9).

    Arguments:
    X (np.ndarray): Feature matrix containing pixel values.
    y (np.ndarray): Labels corresponding to images.
    """

    plt.figure(figsize=(10, 4))
    unique_classes = np.unique(y) # Get unique class labels

    for i, digit in enumerate(unique_classes):
        index = np.where(y == digit)[0][0] # Find first occurrence of the class
        image = X[index].reshape(28, 28) # Reshape 1D array to 28x28

        plt.subplot(2, 5, i + 1)
        plt.imshow(image, cmap='gray')
        plt.title(f"Digit: {digit}")
        plt.axis('off')

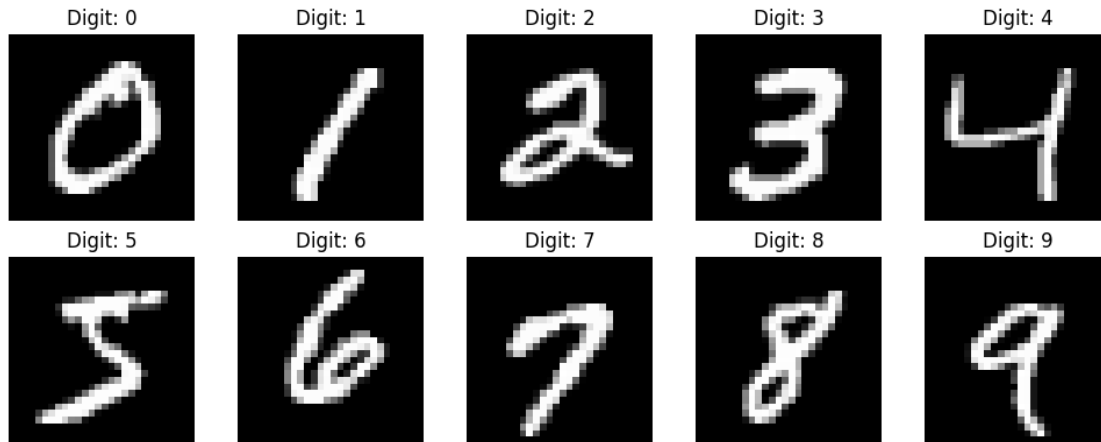
    plt.tight_layout()
    plt.show()

```

```

[22]: csv_file_path = "/content/drive/MyDrive/al ML/mnist_dataset.csv" # Path to
↳saved dataset
X_train, X_test, y_train, y_test = load_and_prepare_mnist(csv_file_path)

```



### 0.3.1 A Quick debugging Step:

```
[23]: # Assert that X and y have matching lengths
assert len(X_train) == len(y_train), f"Error: X and y have different lengths!␣
↪X={len(X_train)}, y={len(y_train)}"
print("Move forward: Dimension of Feture Matrix X and label vector y matched.")
```

Move forward: Dimension of Feture Matrix X and label vector y matched.

```
[20]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

## 0.4 Train the Model:

```
[24]: print(f"Training data shape: {X_train.shape}")
print(f"Test data shape: {X_test.shape}")
```

Training data shape: (48000, 784)

Test data shape: (12000, 784)

```
[25]: from sklearn.preprocessing import OneHotEncoder

# Check if y_train is one-hot encoded
if len(y_train.shape) == 1:
    encoder = OneHotEncoder(sparse_output=False) # Use sparse_output=False for␣
↪newer versions of sklearn
    y_train = encoder.fit_transform(y_train.reshape(-1, 1)) # One-hot encode␣
↪labels
```

```

    y_test = encoder.transform(y_test.reshape(-1, 1)) # One-hot encode test
    ↪ labels

# Now y_train is one-hot encoded, and we can proceed to use it
d = X_train.shape[1] # Number of features (columns in X_train)
c = y_train.shape[1] # Number of classes (columns in y_train after one-hot
    ↪ encoding)

# Initialize weights with small random values and biases with zeros
W = np.random.randn(d, c) * 0.01 # Small random weights initialized
b = np.zeros(c) # Bias initialized to 0

# Set hyperparameters for gradient descent
alpha = 0.1 # Learning rate
n_iter = 1000 # Number of iterations to run gradient descent

# Train the model using gradient descent
W_opt, b_opt, cost_history = gradient_descent_softmax(X_train, y_train, W, b,
    ↪ alpha, n_iter, show_cost=True)


# Plot the cost history to visualize the convergence
plt.plot(cost_history)
plt.title('Cost Function vs. Iterations')
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.grid(True)
plt.show()

```

```

Iteration 1/1000, Cost: 2.209300
Iteration 101/1000, Cost: 0.607304
Iteration 201/1000, Cost: 0.489532
Iteration 301/1000, Cost: 0.440923
Iteration 401/1000, Cost: 0.412845
Iteration 501/1000, Cost: 0.393977
Iteration 601/1000, Cost: 0.380154
Iteration 701/1000, Cost: 0.369446
Iteration 801/1000, Cost: 0.360820
Iteration 901/1000, Cost: 0.353669
Iteration 1000/1000, Cost: 0.347664

```



## 0.5 Evaluating the Model:

```
[26]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score

# Evaluation Function
def evaluate_classification(y_true, y_pred):
    """
    Evaluate classification performance using confusion matrix, precision, recall, and F1-score.

    Parameters:
    y_true (numpy.ndarray): True labels
    y_pred (numpy.ndarray): Predicted labels

    Returns:
    tuple: Confusion matrix, precision, recall, F1 score
```

```

"""
# Compute confusion matrix
cm = confusion_matrix(y_true, y_pred)

# Compute precision, recall, and F1-score
precision = precision_score(y_true, y_pred, average='weighted')
recall = recall_score(y_true, y_pred, average='weighted')
f1 = f1_score(y_true, y_pred, average='weighted')

return cm, precision, recall, f1

```

```

[28]: # Predict on the test set
y_pred_test = predict_softmax(X_test, W_opt, b_opt)

# Evaluate accuracy
y_test_labels = np.argmax(y_test, axis=1) # True labels in numeric form

# Evaluate the model
cm, precision, recall, f1 = evaluate_classification(y_test_labels, y_pred_test)

# Print the evaluation metrics
print("\nConfusion Matrix:")
print(cm)
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1-Score: {f1:.2f}")

# Visualizing the Confusion Matrix
fig, ax = plt.subplots(figsize=(12, 12))
cax = ax.imshow(cm, cmap='Blues') # Use a color map for better visualization

# Dynamic number of classes
num_classes = cm.shape[0]
ax.set_xticks(range(num_classes))
ax.set_yticks(range(num_classes))
ax.set_xticklabels([f'Predicted {i}' for i in range(num_classes)])
ax.set_yticklabels([f'Actual {i}' for i in range(num_classes)])

# Add labels to each cell in the confusion matrix
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, cm[i, j], ha='center', va='center', color='white' if
            cm[i, j] > np.max(cm) / 2 else 'black')

# Add grid lines and axis labels
ax.grid(False)
plt.title('Confusion Matrix', fontsize=14)

```

```
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('Actual Label', fontsize=12)

# Adjust layout
plt.tight_layout()
plt.colorbar(cax)
plt.show()
```

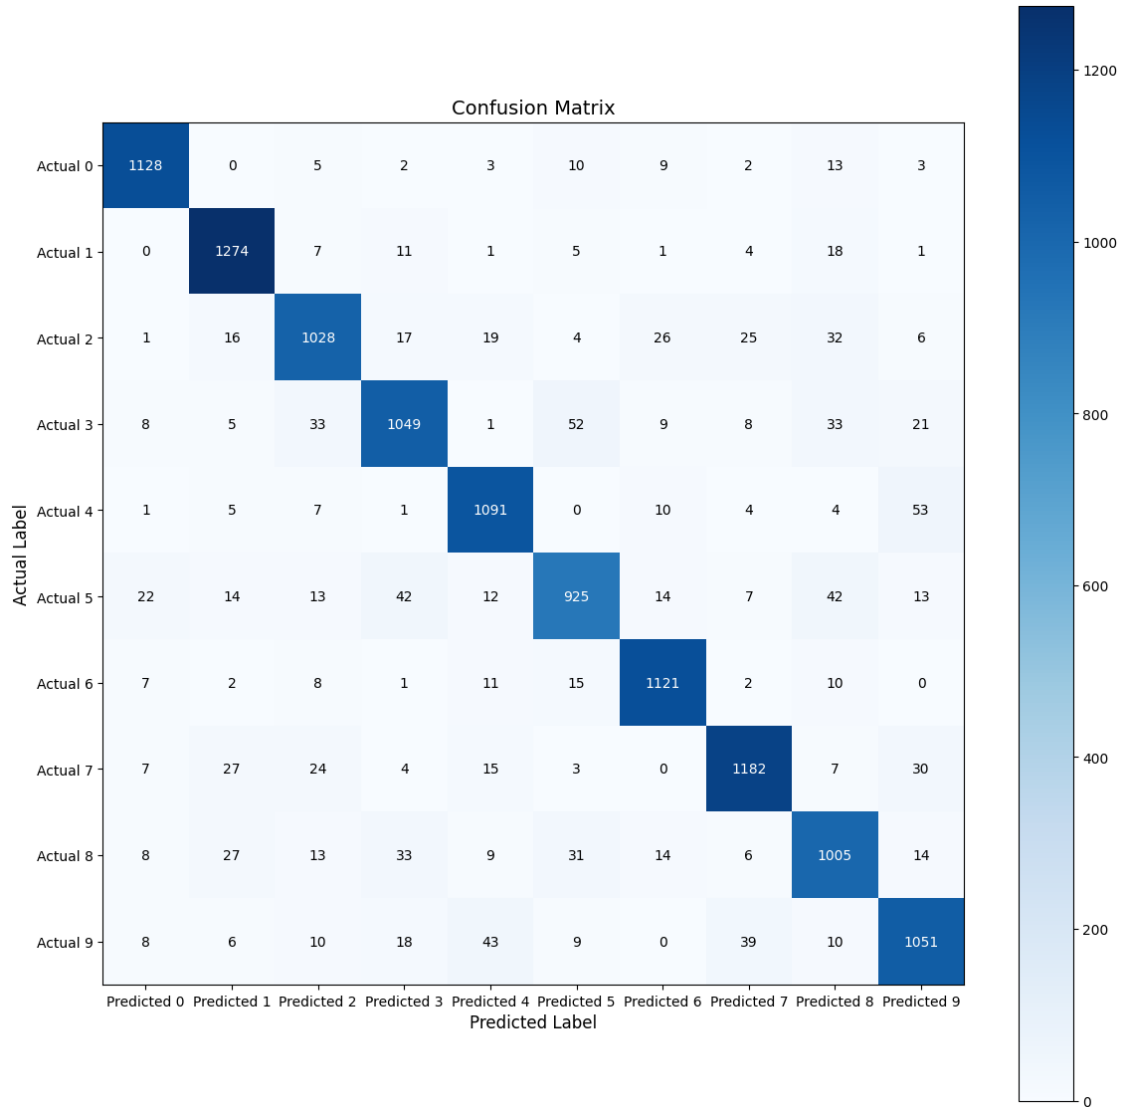
Confusion Matrix:

```
[[1128    0    5    2    3   10    9    2   13    3]
 [    0 1274    7   11    1    5    1    4   18    1]
 [    1   16 1028   17   19    4   26   25   32    6]
 [    8    5   33 1049    1   52    9    8   33   21]
 [    1    5    7    1 1091    0   10    4    4   53]
 [   22   14   13   42   12  925   14    7   42   13]
 [    7    2    8    1   11   15 1121    2   10    0]
 [    7   27   24    4   15    3    0 1182    7   30]
 [    8   27   13   33    9   31   14    6 1005   14]
 [    8    6   10   18   43    9    0   39   10 1051]]
```

Precision: 0.90

Recall: 0.90

F1-Score: 0.90



## 1 Linear Seperability and Logistic Regression:

```
[27]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification, make_circles
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

# Set random seed for reproducibility
np.random.seed(42)

# Generate linearly separable dataset
```

```

X_linear_separable, y_linear_separable = make_classification(n_samples=200,
    ↪n_features=2,
    n_informative=2,
    ↪n_redundant=0,
    ↪n_clusters_per_class=1, random_state=42)

# Split the data into training and testing sets
X_train_linear, X_test_linear, y_train_linear, y_test_linear = train_test_split(
    X_linear_separable, y_linear_separable, test_size=0.2, random_state=42
)

# Train logistic regression model on linearly separable data
logistic_model_linear_separable = LogisticRegression()
logistic_model_linear_separable.fit(X_train_linear, y_train_linear)

# Generate non-linearly separable dataset (circles)
X_non_linear_separable, y_non_linear_separable = make_circles(n_samples=200,
    ↪noise=0.1, factor=0.5,
    random_state=42)

# Split the data into training and testing sets
X_train_non_linear, X_test_non_linear, y_train_non_linear, y_test_non_linear =
    ↪train_test_split(
        X_non_linear_separable, y_non_linear_separable, test_size=0.2,
        ↪random_state=42
    )

# Train logistic regression model on non-linearly separable data
logistic_model_non_linear_separable = LogisticRegression()
logistic_model_non_linear_separable.fit(X_train_non_linear, y_train_non_linear)

# Plot decision boundaries for linearly and non-linearly separable data
def plot_decision_boundary(ax, model, X, y, title):
    h = .02 # step size in the mesh
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    ax.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
    ax.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Paired)
    ax.set_title(title)
    ax.set_xlabel('Feature 1')
    ax.set_ylabel('Feature 2')

```



```

# Create subplots
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

# Plot decision boundary for linearly separable data (Training)
plot_decision_boundary(axes[0, 0], logistic_model_linear_separable,
    ↪X_train_linear, y_train_linear,
    'Linearly Separable Data (Training)')

# Plot decision boundary for linearly separable data (Testing)
plot_decision_boundary(axes[0, 1], logistic_model_linear_separable,
    ↪X_test_linear, y_test_linear,
    'Linearly Separable Data (Testing)')

# Plot decision boundary for non-linearly separable data (Training)
plot_decision_boundary(axes[1, 0], logistic_model_non_linear_separable,
    ↪X_train_non_linear,
    y_train_non_linear, 'Non-Linearly Separable Data
    ↪(Training)')

# Plot decision boundary for non-linearly separable data (Testing)
plot_decision_boundary(axes[1, 1], logistic_model_non_linear_separable,
    ↪X_test_non_linear,
    y_test_non_linear, 'Non-Linearly Separable Data
    ↪(Testing)')

plt.tight_layout()

# Save the plots as PNG files
plt.savefig('decision_boundaries.png')
plt.show()

```

